

INF4820, Algorithms for AI and NLP:

Evaluating Classifiers
Clustering

Erik Velldal

University of Oslo

Sept. 18, 2012



Classification

- ▶ Recap
- ▶ Evaluating classifiers
 - ▶ Accuracy, precision, recall and F-score

Clustering

- ▶ Unsupervised machine learning for class discovery.
- ▶ Flat vs. hierarchical clustering.
- ▶ Example of flat / partional clustering: k -means clustering.




- ▶ Supervised vs unsupervised learning.
- ▶ Vectors space classification.
- ▶ How to represent classes and class membership.
- ▶ Rocchio + k NN.
- ▶ Linear vs non-linear decision boundaries.



- ▶ We've seen how vector space classification amounts to computing the boundaries in the space that separate the class regions;
the decision boundaries.
- ▶ To evaluate the boundary, we measure the number of correct classification predictions on unseen test items.
 - ▶ Many ways to do this. . .
- ▶ Why can't we test on the training data?
- ▶ We want to test how well a model *generalizes* on a **held-out** test set.
- ▶ (Or, if we have little data, by n -fold cross-validation.)
- ▶ Labeled test data is sometimes referred to as the **gold standard**.

Example: Evaluating classifier decisions




- Predictions for a given class can be wrong or correct in two ways:

	gold = positive	gold = negative
prediction = positive	true positive (TP)	false positive (FP)
prediction = negative	false negative (FN)	true negative (TN)



- ▶ *accuracy* = $\frac{TP+TN}{N} = \frac{TP+TN}{TP+TN+FP+FN}$
 - ▶ The ratio of correct predictions.
 - ▶ Not suitable for unbalanced numbers of positive / negative examples.
- ▶ *precision* = $\frac{TP}{TP+FP}$
 - ▶ The number of detected class members that were correct.
- ▶ *recall* = $\frac{TP}{TP+FN}$
 - ▶ The number of actual class members that were detected.
 - ▶ Trade-off: Positive predictions for all examples would give 100% recall but (typically) terrible precision.
- ▶ *F-score* = $\frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$
 - ▶ Balanced measure of precision and recall (harmonic mean).

Example: Evaluating classifier decisions



$$\begin{aligned} \text{accuracy} &= \frac{TP+TN}{N} \\ &= \frac{1+6}{10} = 0.7 \end{aligned}$$

$$\begin{aligned} \text{precision} &= \frac{TP}{TP+FP} \\ &= \frac{1}{1+1} = 0.5 \end{aligned}$$

$$\begin{aligned} \text{recall} &= \frac{TP}{TP+FN} \\ &= \frac{1}{1+2} = 0.33 \end{aligned}$$

$$\begin{aligned} F\text{-score} &= \\ \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}} &= 0.4 \end{aligned}$$



Macro-averaging

- ▶ Sum precision and recall for each class, and then compute global averages of these.
- ▶ The **macro** average will be highly influenced by the **small** classes.

Micro-averaging

- ▶ Sum TPs, FPs, and FNs for all points/objects across all classes, and then compute global precision and recall.
- ▶ The **micro** average will be highly influenced by the **large** classes.



Classification

- ▶ **Supervised** learning, requiring **labeled** training data.
- ▶ Given some training set of examples with class labels, train a classifier to predict the class labels of unseen objects.



Clustering

- ▶ **Unsupervised** learning from **unlabeled** data.
- ▶ Automatically group similar objects together.
- ▶ No pre-defined classes: we only specify the similarity measure.
- ▶ Described by ? (?) as "*the search for structure in data*".
- ▶ General objective:
 - ▶ Partition the data into subsets, so that the similarity among members of the same group is high (**homogeneity**) while the similarity between the groups themselves is low (**heterogeneity**).



- ▶ Visualization and exploratory data analysis.
- ▶ Generalization and abstraction. “Reason by analogy”.
 - ▶ Can have class-based models, even without predefined classes.
 - ▶ Helps alleviating the sparse data problem.
- ▶ Many applications within IR. Examples:
 - ▶ Speed up search: First retrieve the most relevant cluster, then retrieve documents from within the cluster.
 - ▶ Presenting the search results: Instead of ranked lists, organize the results as clusters (see e.g. clusty.com).
- ▶ News aggregation / topic directories.
- ▶ Social network analysis; identify sub-communities and user segments.
- ▶ Image segmentation, product recommendations, demographic analysis, ...

Different methods can be divided according to the *memberships* they create and the *procedure* by which the clusters are formed:





Hierarchical

- ▶ Creates a tree structure of hierarchically nested clusters
- ▶ **Divisive** (top-down): Let all objects be members of the same cluster; then successively split the group into smaller and maximally dissimilar clusters until all objects is its own singleton cluster.
- ▶ **Agglomerative** (bottom-up): Let each object define its own cluster; then successively merge most similar clusters until only one remains.

Flat

- ▶ Often referred to as **partitional clustering** when assuming hard and disjoint clusters. (But can also be soft.)
- ▶ Tries to directly decompose the data into a set of clusters.



- ▶ Given a set of objects $O = \{o_1, \dots, o_n\}$, a hard flat clustering algorithm seeks to construct a set of clusters $C = \{c_1, \dots, c_k\}$, where each object o_i is assigned to a single cluster c_i .
- ▶ The **cardinality** k (= the number of clusters) must typically be manually specified as a parameter to the algorithm.
- ▶ But the most important parameter is the **similarity function** s .
- ▶ More formally, we want to define an assignment $\gamma : O \rightarrow C$ that optimizes some objective function $F_s(\gamma)$.
- ▶ The objective function is defined in terms of the similarity function, and generally we want to optimize for:
 - ▶ High intra-cluster similarity
 - ▶ Low inter-cluster similarity



Optimization problems are search problems:

- ▶ There's a finite number of possible partitionings of O .
- ▶ Naive solution: enumerate all possible assignments $\Gamma = \{\gamma_1, \dots, \gamma_m\}$ and choose the best one,

$$\hat{\gamma} = \arg \min_{\gamma \in \Gamma} F_s(\gamma)$$

- ▶ Problem: Exponentially many possible partitions.
- ▶ Approximate the solution by iteratively improving on an initial (possibly random) partition until some stopping criterion is met.



- ▶ Unsupervised variant of the Rocchio classifier.
- ▶ **Goal:** Partition the n observed objects into k clusters C so that each point \vec{x}_j belongs to the cluster c_i with the nearest centroid $\vec{\mu}_i$.
- ▶ Typically assumes Euclidean distance as the similarity function s .
- ▶ **The optimization problem:** For each cluster, minimize the *within-cluster sum of squares*, $F_s = \text{WCSS}$:

$$\text{WCSS} = \sum_{c_i \in C} \sum_{\vec{x}_j \in c_i} \|\vec{x}_j - \vec{\mu}_i\|^2$$

- ▶ WCSS also amounts to the more general measure of how well a model fits the data known as the *residual sum of squares* (RSS).
- ▶ Minimizing RSS is equivalent to minimizing the average squared distance between objects and their cluster centroids (since n is fixed), —a measure of how well each centroid represents the members assigned to the cluster.



Algorithm

Initialize: Compute centroids for k seeds.

Iterate:


- Assign each object to the cluster with the nearest centroid.
- Compute new centroids for the clusters.

Terminate: When stopping criterion is satisfied.


Properties

- ▶ In short, we iteratively reassign memberships and recompute centroids until the configuration stabilizes.
- ▶ WCSS is monotonically decreasing (or unchanged) for each iteration.
- ▶ Guaranteed to converge but not to find the global minimum.
- ▶ The time complexity is linear, $O(kn)$.


k -Means example for $k = 2$ in R^2 (Manning, Raghavan & Schütze 2008)




selection of seeds




assignment of documents (iter. 1)



recomputation/movement of $\bar{\mu}$'s (iter. 1)



$\bar{\mu}$'s after convergence (iter. 9)



movement of $\bar{\mu}$'s in 9 iterations



"Seeding"

- ▶ We initialize the algorithm by choosing random *seeds* that we use to compute the first set of centroids.
- ▶ Many possible heuristics for selecting the seeds:
 - ▶ pick k random objects from the collection;
 - ▶ pick k random points in the space;
 - ▶ pick k sets of m random points and compute centroids for each set;
 - ▶ compute an hierarchical clustering on a subset of the data to find k initial clusters; etc..
- ▶ The initial seeds can have a large impact on the resulting clustering (because we typically end up only finding a local minimum of the objective function).
- ▶ **Outliers** are troublemakers.



Possible termination criterions

- ▶ Fixed number of iterations
- ▶ Clusters or centroids are unchanged between iterations.
- ▶ Threshold on the decrease of the objective function (absolute or relative to previous iteration)

Some Close Relatives of k -Means

- ▶ **k -Medoids**: Like k -means but uses medoids instead of centroids to represent the cluster centers.
- ▶ **Fuzzy c -Means** (FCM): Like k -means but assigns soft memberships in $[0, 1]$, where membership is a function of the centroid distance.
 - ▶ The computations of both WCSS and centroids are weighted by the membership function.



Pros

- ▶ Conceptually simple, and easy to implement.
- ▶ Efficient. Typically linear in the number of objects.

Cons

- ▶ The dependence on the random seeds makes the clustering **non-deterministic**.
- ▶ The number of clusters k must be pre-specified. Often no principled means of *a priori* specifying k .
- ▶ The clustering quality often considered inferior to that of the less efficient hierarchical methods.
- ▶ Not as informative as the more structured clusterings produced by hierarchical methods.



- ▶ Agglomerative vs. divisive hierarchical clustering
- ▶ Reading: Chapter 17 in Manning, Raghavan & Schütze (2008), *Introduction to Information Retrieval*;
<http://informationretrieval.org/>
(see course web-page for the relevant sections).